

CS 2302 Data Structures

Fall 2017

B-Trees

1 Definition

A B-tree is a balanced search tree with the following properties:

- Each node x has the following fields:
 - $x.n$ the number of keys stored in x
 - $x.key[i]$, $0 \leq i < x.n$ the keys
 - $x.isLeaf$ a Boolean indicator that is true if x is a leaf
 - $x.c[i]$, $0 \leq i \leq x.n$, an array of references to the children of x
- The keys separate the ranges of keys stored in every subtree of x . if k_i is stored in a subtree with root $x.c[i]$, then $k_0 < x.key[0] < k_1 < x.key[1] < x.key[x.n - 1] < k_{x.n}$
- Every leaf has the same depth, which is the tree's height
- Every B-tree has a parameter t called the *minimum degree*. Every node, with the possible exception of the root, must have at least $t - 1$ keys (and t children if it is not a leaf). Every node must have at most $2t - 1$ keys (and $2t$ children if it is not a leaf)

1.1 Building a B-tree:

Building a B-tree:

- Insertions occur at leaves.
- We start at the root and descend until we find the leaf where the new element should be stored.
- If a full node is found (one that contains $2t - 1$ keys, or, equivalently, where $x.n == x.key.length$) while tracing the path to the appropriate leaf to insert the new element, split it, adding median key to its parent.
- The height of the tree is increased when we split the root.

As a consequence of the building process, all leaves have the same depth, which is the height of the tree.

1.2 Class definition:

```
public class BTreeNode{
    public int n;           // Actual number of keys on the node
    public boolean isLeaf;  // Boolean indicator, true for leaf nodes, false otherwise
    public int[] key;       // Keys stored in the node. They are sorted in ascending order
    public BTreeNode[] c;   // Children of node. Keys in c[i] are less than key[i] (if it exists)
                           // and greater than key[i-1] if it exists

    public BTreeNode(int t){ // Build empty node
        isLeaf = true;
        key = new int[2*t-1]; // Array sizes are set to maximum possible value
        c = new BTreeNode[2*t];
        n=0;                  // Number of elements is zero, since node is empty
    }
}
```

2 Problem Solving

We'll cover many different problems using B-trees; however, they can be classified into three types, each involving very similar operations. Understanding how each of them works is crucial to develop your problem-solving skills using this data structure.

2.1 Basic problem types:

Ordered by increasing difficulty, the three problem types are:

1. Traverse a predefined path on the tree. Examples: finding the height of a tree, finding the maximum or minimum key in the tree.
2. Traverse the whole tree. Examples: printing all the keys in the tree, counting the number of keys in the tree, counting the number of nodes in the tree.
3. Searching for a given key. This requires following a path that depends on the given key and the contents of the tree.

2.2 Examples:

Problem type I: finding the smallest key in the tree.

```
int minimum(BTreeNode T){
    if (T.isLeaf)
        return T.key[0];
    return minimum(T.c[0]);
}
```

Problem type II: printing the keys in a B-tree in ascending order.

```
public void printKeys(BTreeNode T){
    if (T.isLeaf){
        for(int i =0; i<T.n;i++)
            System.out.print(T.key[i]+" ");
    }
    else{
        for(int i =0; i<T.n;i++){
            printKeys(T.c[i]);
            System.out.print(T.key[i]+" ");
        }
        printKeys(T.c[T.n]);
    }
}
```

Problem type III: finding the node that contains a given key k given a reference to the root of the tree.

```
BTreeNode SearchBTree(BTreeNode T, int k){
    int i=0;
    while ((i<T.n )&&(k>T.key[i])) //Sequentially search for k
        i++;
    if (i==T.n) || (k<T.key[i]) //k is not in current node.
        if T.isleaf
            return null;
        else
            return SearchBTree(T.c[i], k);
    else
        return T;           //k is in current node.
}
```

2.3 Dealing with depth:

For each of the three types of problems, we may receive an extra parameter or return a value that indicates the depth of the node(s) we are interested in.

2.3.1 Examples:

Problem type I: finding the largest key in the tree at a given depth d . If the height of the tree is less than d return Integer.MIN_VALUE.

```
int maximumAtDepthD(BTreeNode T, int d){
    if (d==0)
        return T.key[T.n-1];
}
```

```

    if (T.isLeaf)
        return Integer.MIN_VALUE;
    return maximumAtDepthD(T.c[T.n],d-1);
}

```

Problem type II: printing the keys at depth d in a B-tree in ascending order.

```

public void printKeys(BTreeNode T, int d){
    if (d==0)
        for(int i =0; i<T.n;i++)
            System.out.print(T.key[i]+" ");
    else
        if(!T.isLeaf)
            for(int i =0; i<=T.n;i++)
                printKeys(T.c[i],d-1);
}

```

Problem type III: return the depth at which a given key k is found or -1 if k is not in the tree.

```

int SearchBTree(BTreeNode T, int k){
    int i=0;
    while ((i<T.n )&&(k>T.key[i]))//Sequentially search for k
        i++;
    if (i==T.n) || (k<T.key[i]) // k is not in current node.
        if (T.isleaf) // k is not in the tree
            return -1;
        else{
            int d = SearchBTree(T.c[i], k);
            if (d==-1) // k is not in the sub-tree
                return -1;
            else
                return d+1;
        }
    else
        return 0; //k is in current node.
}

```